

# SEARCH – a New Slow Start Algorithm for TCP and QUIC


Jae Chung  
Feng Li

Maryam Ataei Kachooei  
Mark Claypool


*IETF CCWG*  
Vancouver, Canada  
July 2024




# Motivation



Exit Too Early




Exit Too Late




Exit at Choke Point

TCP HyStart over Wireless?


**GEO**




**4G LTE**



**LEO**



**Wi-Fi**







# Outline

- Motivation (done)
- SEARCH (next)
- Performance Evaluation
- Conclusion



# SEARCH – Slow start Exit at Right CHokepoint





# SEARCH – Slow start Exit at Right CHokepoint


$$\text{sent}' = 2 \cdot \text{delv}_{\text{previous}}$$

$$\text{diff} = \text{sent}' - \text{delv}_{\text{now}}$$

$$\text{normalized\_diff} = \text{diff} / \text{sent}'$$

$$\text{normalized\_diff} \geq \text{threshold?}$$

→ exit slow start






# Challenges


## Variable RTTs

- Caused by uplink ACK schedule timing
- *Not* caused by congestion on forward link



## Limit memory on server

- Memory allocated per flow
- Can't store history for each ACK





# SEARCH

```
Parameters:
0: WINDOW_FACTOR = 3.5
1: W = 10
2: EXTRA_BINS = 15
3: NUM_BINS = W + EXTRA_BINS
4: THRESH = 0.35

Initialization():
5: window_size = initial_rtt x WINDOW_FACTOR
6: bin_duration = window_size / W
7: bin[NUM_BINS] = {}
8: curr_idx = -1
9: bin_end = now + bin_duration

ACK_arrived(sequence_num, rtt):

    // Check if passed bin boundary.
10: if (*now* > bin_end) then
11:     update_bins()

    // Check if enough data for SEARCH.
12: prev_idx = curr_idx - (rtt / bin_duration)
13: if (prev_idx >= W) and
14:     (curr_idx - prev_idx) <= EXTRA_BINS then

        // Run SEARCH check.
15: curr_delv = compute_delv(curr_idx - W, curr_idx)
16: fraction = (rtt mod bin_duration) / bin_duration
17: prev_delv = compute_delv(prev_idx - W, prev_idx, fraction)
18: norm_diff = (2 x prev_delv - curr_delv) / (2 x prev_delv)
19: if (norm_diff >= THRESH) then
20:     exit_slow_start()
21: end if

22: end if // Enough data for SEARCH.
23: end if // Each ACK.

// Update bin statistics, accounting for cases where more
// than one bin boundary might have been passed.
update_bins():
24: passed_bins = (*now* - bin_end) / bin_duration + 1
25: bin_end += passed_bins x bin_duration
26: for i = (curr_idx + 1) to (curr_idx + passed_bins)
27:     if (curr_idx >= 0) bin[i mod NUM_BINS] = bin[curr_idx]
28: end for
29: curr_idx += passed_bins
30: bin[curr_idx mod NUM_BINS] = sequence_num

// Compute delivered bytes over the window of bins, interpolating a
// fraction of each bin on the end (default is 0).
compute_delv(idx1, idx2, fraction = 0):
31: delv = 0
32: delv = bin[(idx2 - 1) mod NUM_BINS] - bin[idx1 mod NUM_BINS]
33: delv += (bin[idx1 mod NUM_BINS] - bin[(idx1 - 1) mod NUM_BINS]) x fraction
34: delv += (bin[idx2 mod NUM_BINS] - bin[(idx2 - 1) mod NUM_BINS]) x (1 - fraction)
35: return delv

// Exit slow start by setting cwnd and ssthresh.
exit_slow_start():
36: cong_idx = curr_idx - 2 x initial_rtt / bin_duration
37: overshoot = compute_delv(cong_idx, curr_idx)
38: cwnd -= overshoot
39: ssthresh = cwnd
```

Initialize

Receive  
ACK

Curr, Prev,  
Norm diff

Norm diff >  
THRESH?

Update  
bins

Exit slow  
start





# Parameter Selection

- **Window size:**
  - Large enough for link RTT variation
  - Small enough to respond quickly
- **Number of bins:**
  - Large enough to reduce load
  - Small enough to maintain fidelity and respond quickly
- **Threshold:**
  - Large enough so above noise
  - Small enough to respond quickly

## Improving TCP Slow Start Performance in Wireless Networks with SEARCH

*IEEE World of Wireless, Mobile and Multimedia Networks (WoWMoM)*

Perth, Australia, June 2024.

0: WINDOW\_FACTOR = **3.5**

1: W = **10**

2: EXTRA\_BINS = **15**

3: NUM\_BINS = W + EXTRA\_BINS

4: THRESHOLD = **0.35**



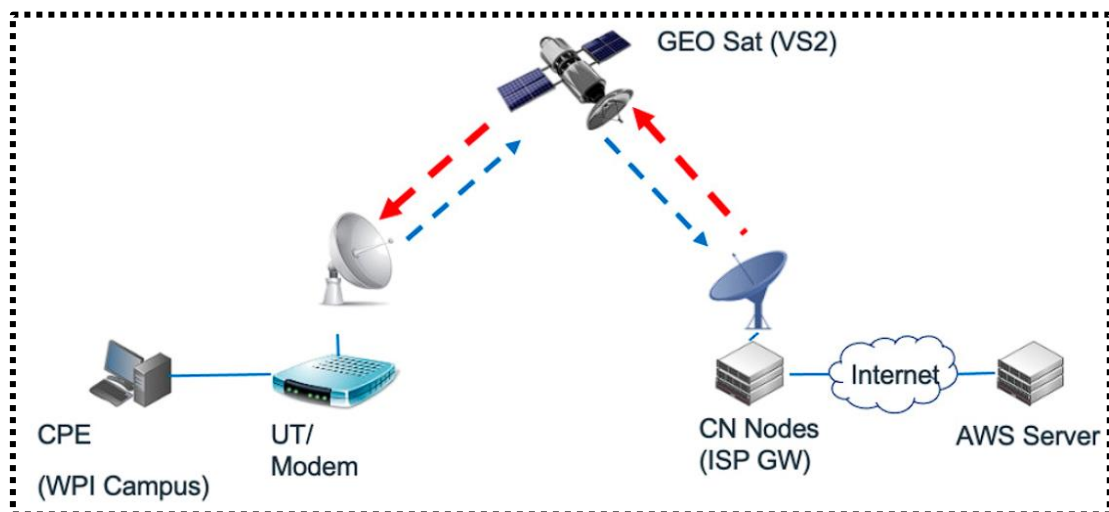
# Outline

- Motivation
- SEARCH
- Performance Evaluation

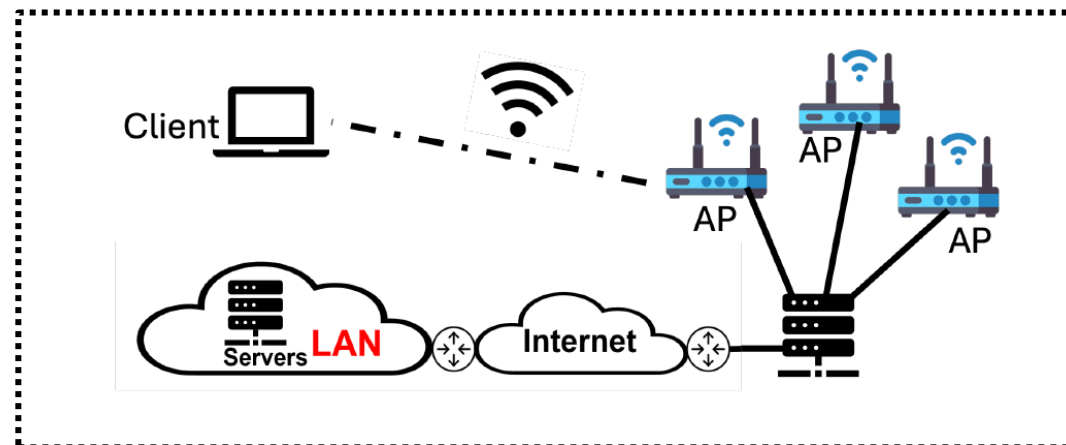
(done)

(done)

(next)



**GEO**



**Wi-Fi**

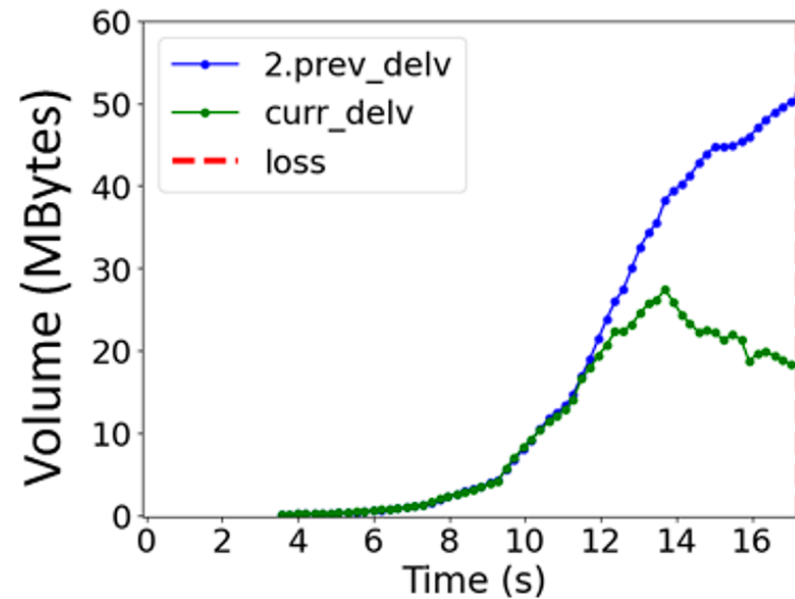
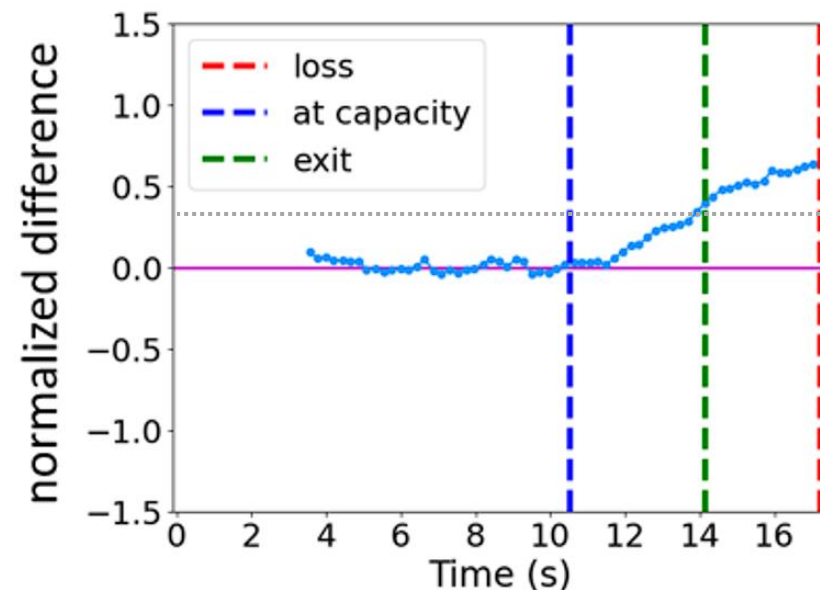






# Performance – GEO Satellite



SEARCH

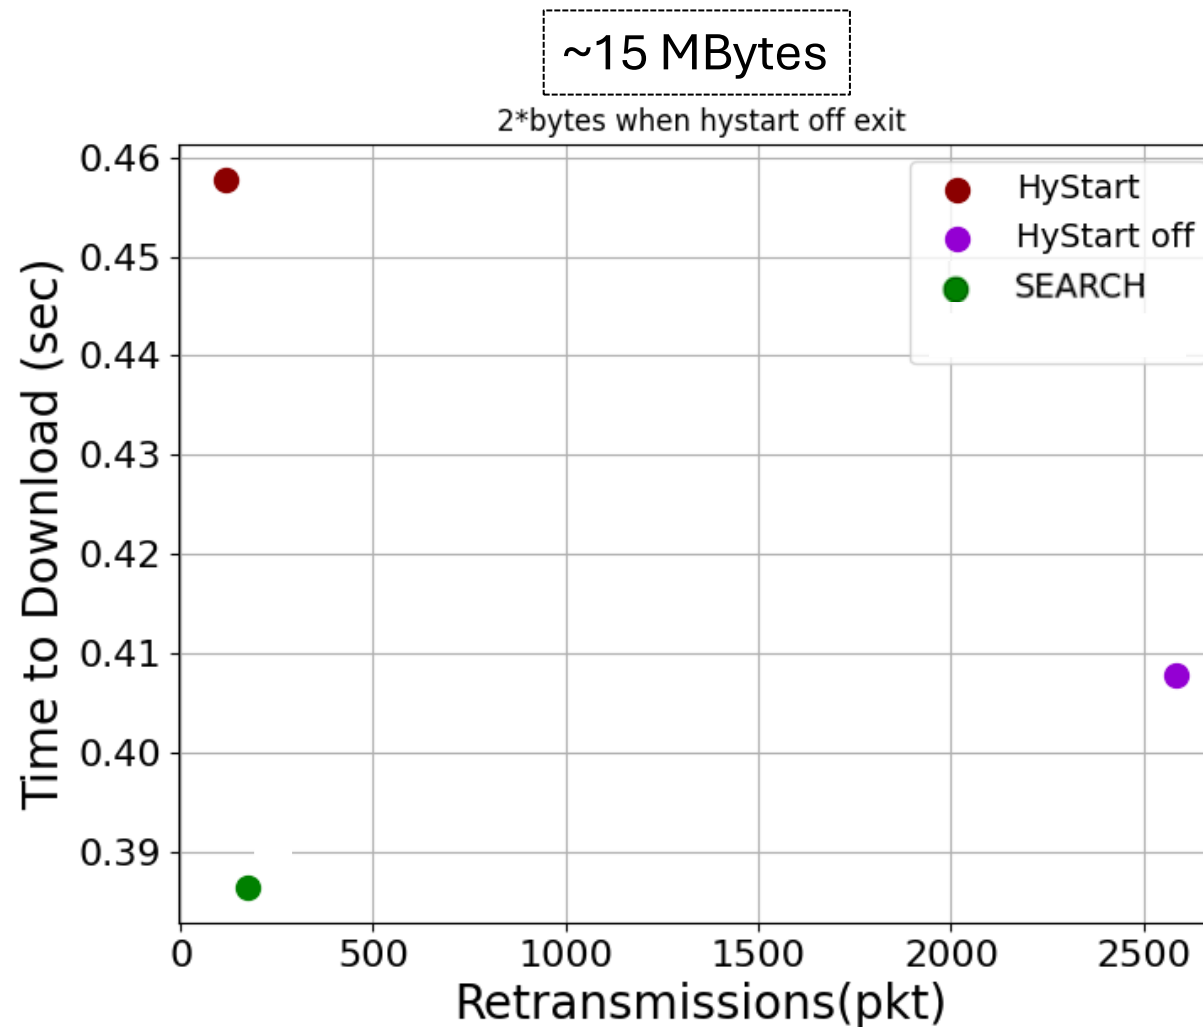

HyStart Off

HyStart On





# Performance – Wi-Fi





# Implementation Status

- Linux Kernel Modules

- 5.13.x series kernel (main branch)

- [https://github.com/Project-Faster/tcp\\_ss\\_search.git](https://github.com/Project-Faster/tcp_ss_search.git)

- 6.10rc2 based (net-next-6.10rc2 branch)

- [https://github.com/Project-Faster/tcp\\_ss\\_search/tree/net-next-6.10rc2](https://github.com/Project-Faster/tcp_ss_search/tree/net-next-6.10rc2)

- QUIC H2O/Quicly

- <https://github.com/Project-Faster/quicly/tree/generic-slowstart>

- Upstream into Linux mainstream and open source QUIC



# Conclusion

- HyStart does not work in wireless environments (GEO, LEO, 4G LTE, Wi-Fi) → premature slow start exits
- **SEARCH**
  - Determines “choke point” from expected delivered bytes
  - Exits slow start after congestion point, before loss
  - Improves utilization, reduces packet loss (versus off)



**SEARCH**

<https://search-ss.wpi.edu/>

Thank-you for your attention!

# SEARCH – a New Slow Start Algorithm for TCP and QUIC

Jae Chung

Feng Li

Maryam Ataei Kachooei

Mark Claypool

*IETF CCWG*

Vancouver, Canada

July 2024





# References

- Improving TCP Slow Start Performance in Wireless Networks with SEARCH
  - *IEEE World of Wireless, Mobile and Multimedia Networks (WoWMoM)*
  - Perth, Australia, June 2024.
- Improving QUIC Slow Start Behavior in Wireless Networks with SEARCH
  - *IEEE Local and Metropolitan Area Networks (LANMAN)*
  - Boston, Massachusetts, USA, July 2024
- Implementation of the SEARCH Slow Start Algorithm in the Linux Kernel
  - *0x18 NetDev Conference*
  - Santa Clara, California, USA, July 2024