

# BBR Congestion Control Draft

## draft-ietf-ccwg-bbr-02

Internet Draft Editors:

Neal Cardwell (Google), Ian Swett (Google), Joseph Beshay (Meta)

Speaker: Ian Swett

# Outline

## Overview

- Outline recent BBR Internet Draft updates
- Summarize open pull requests and open issues

## Goals for this talk:

- Provide a road map for...
  - Readers of the draft
  - Implementers of BBR reading the draft
  - Members of the CCWG/ICCRG community who would like to contribute
- Inviting the community to...
  - Read the draft
  - Contribute to the draft

# Overview of draft-ietf-ccwg-bbr

- BBR was [adopted a CCWG WG item in October 2024](#)
- Intended status: experimental RFC
- IETF CCWG members are collaborating on github:
  - <https://github.com/ietf-wg-ccwg/draft-ietf-ccwg-bbr>
- Latest published revision is at:
  - <https://datatracker.ietf.org/doc/draft-ietf-ccwg-bbr/>
- Latest editor's copy (with hot-off-the-press changes not in published revisions above):
  - <https://ietf-wg-ccwg.github.io/draft-ietf-ccwg-bbr/draft-ietf-ccwg-bbr.html>
- Draft editors:
  - Neal Cardwell (Google)
  - Ian Swett (Google)
  - Joseph Beshay (Meta)

# Changes in draft-ietf-ccwg-bbr-02: summary

- Previous version was: draft-ietf-ccwg-bbr-01, discussed at IETF 121 in Nov, 2024 [[slides](#)]
- Diffs: [[text diff between draft-ietf-ccwg-bbr-01 and draft-ietf-ccwg-bbr-02](#)]
- Main changes between draft-ietf-ccwg-bbr-01 and draft-ietf-ccwg-bbr-02:
  - Add missing BBRIsProbingBW() pseudocode [#10](#)
  - Ensure BBRCheckFullBWReached() pseudocode counts three full rounds [#16](#)
  - Remove some TCP and SACK references [#15](#)
  - Rename *inflight\_lo* and *inflight\_hi* to *inflight\_shortterm* and *inflight\_longterm* [#11](#)
  - Define "acknowledged" and "delivered" in a transport-agnostic way [#21](#)
  - ... a number of other editorial improvements

# Goals of evolving the BBR draft text

Goals as we evolve the BBR draft text:

- Clarification
- Simplification
- Better coexistence with Reno/CUBIC
- Better performance
- Avoiding performance regressions in the real world

Proposed bar for publication (keep in mind the target is an experimental RFC):

- Multiple deployments at scale in QUIC and TCP
- Text both TCP and QUIC implementations can follow
- Fair sharing with other BBR flows, coexistence with Reno and Cubic

Thesis: It's better to publish a good draft with deployment experience in a reasonable timeframe than evolve BBR indefinitely without shipping an RFC.

## 2 Open Pull Requests: minor algorithm changes

2 [Open PRs](#); both are algorithm changes waiting for performance data from experiments:

- [#6 Use consistent value for drain pacing gain which matches derivation doc](#)
  - Proposes changing BBR DRAIN gain from  $1/2.89 = 0.35$  to  $1/2 = 0.5$
  - To match [analytical derivation of DRAIN pacing gain](#), which derives  $1/2 = 0.5$
  - Have some old A/B experiment data from Linux TCP YouTube experiments:
    - Unclear if there are statistically significant performance regressions
    - Planning to re-run an experiment to ensure there are no regressions
  - We'd appreciate performance data comparing drain\_gain vals from anyone who's able
- [#5 Remove BBR.ack\\_phase from pseudocode](#)
  - A minor algorithm simplification
  - Has one implementation (mvfst QUIC BBR2)
  - We'd like
    - A second implementation
    - Internet performance data to verify there is no performance regression

# Issue: generalization to non-TCP transports

Open issue:

- Bandwidth estimation still uses seq names, which are TCP-centric

The intent is to make the draft as transport agnostic as possible.

We are making progress but not done.

Goal: Ensure implementation of BBR across as many transports as possible

Non-Goal: Create universal approach for mapping any congestion control to any transport

# Issue: ProbeRTT every 5 sec or 10 sec?

Open issue:

- [Impact of increasing ProbeRTT interval to 10s?](#)

Details:

- BBRv1 used a 10s ProbeRTT interval
  - This 10s value has many "air miles"
- BBRv2/BBRv3 currently use: "ProbeRTTInterval = 5 seconds"
  - Rationale: since BBRv2/BBRv3 ProbeRTT reduce cwnd to  $0.5 * \text{estimated\_BDP}$  instead of 4 packets, they can achieve the same throughput as BBRv1 with 2x more frequent ProbeRTT
- Meta mvfst QUIC BBRv2 found:
  - Increasing ProbeRTT interval to 10s in internet facing connections improves application performance in initial experiments (vs 5s)
    - No increase in rtt estimates
    - No increase in retransmissions
    - Reduces p99 response time for requests by up to 15% for app traffic on busy links



# Issue: "inflight" or "in-flight data" or "data in flight"

Open issue:

- "inflight" or "in-flight data" or "data in flight"
- All 3 phrases are used in the draft today
  - RFC 9002 (QUIC) almost always uses 'in flight'
- We should probably standardize on one
- Purely editorial

# Issue: Do we need inflight\_shortterm?

Open issue:

- [Do we need inflight\\_shortterm?](#)

Propose resolving this with "Yes"; rationale ([here](#)):

- When available bandwidth is reduced (e.g., due to more flows competing at the bottleneck or a reduced bottleneck link rate, e.g. radio issues) this reduces not just the bandwidth available but also the flow's fair share of the path BDP and bottleneck buffer.
- Thus, to avoid excess queuing and loss, the flow needs to reduce its pacing rate so that matches the lower available bandwidth (via bw\_shortterm), and reduce its cwnd to match the lower share of the path BDP and bottleneck buffer slots (via inflight\_shortterm).

Notes:

- If folks have simulation tests, lab tests, or production data showing that good results are achievable w/o inflight\_lo (inflight\_shortterm) or equivalent, let's re-open

# Issue: decide/document draft's stance on ECN

Open issue:

- [Decide and document the BBR draft's stance on ECN](#)

Discussion:

- Do we want to document a recommended response to ECN in the draft?
- No strong opinion on the details as long as it performs well in real-world testing
- Not aware of large-scale real-world test results with BBR w/ ECN over public Internet
- Waiting to implement and test an ECN response for BBR may hold up the draft considerably

# Conclusion

- Inviting the community to...
  - Read the draft: [draft-ietf-ccwg-bbr](#)
  - Offer contributions/comments/edits, in whatever manner you prefer
- Thanks!



# Changes in draft-ietf-ccwg-bbr: how to view

- To see recent changes, you can use several approaches, depending on your preference:
  - From the command line:
    - `git clone https://github.com/ietf-wg-ccwg/draft-ietf-ccwg-bbr.git`
    - `cd draft-ietf-ccwg-bbr/`
    - `git log -p`
  - From github:
    - [Commits](#)
    - [Merged pull requests](#)


# Thoughts about ways to contribute

- Contributions at any "rung of the ladders" below are welcome!
- The higher on the "ladders" (the more concrete/specific/tested the contribution is)...
  - The more useful to the BBR draft effort
  - Given editor time constraints, the more likely the eventual inclusion in the draft
- To finalize significant algorithm changes, we'd like to ultimately reach the top rung of the ladder
- Collaboration encouraged: e.g., idea from person A, implemented by person B, tested by sites B/C

## Editorial changes:



- Github pull request with draft text
- Github issue describing the idea
- CCWG email/meeting suggestion



## Technical algorithm changes:



- Multiple at-scale Internet deployments
- At-scale Internet deployment data
- Lab/simulation experiment results
- Patch to an open-source BBR code base
- Github pull request with draft text
- Pseudocode
- Github issue describing the idea
- CCWG email/meeting suggestion