

UPDATE

SEARCH – a New Slow Start Algorithm for TCP and QUIC

Jae Chung
Feng Li

Maryam Ataei Kachooei
Mark Claypool

IETF CCWG
Bangkok, Thailand
March 2025





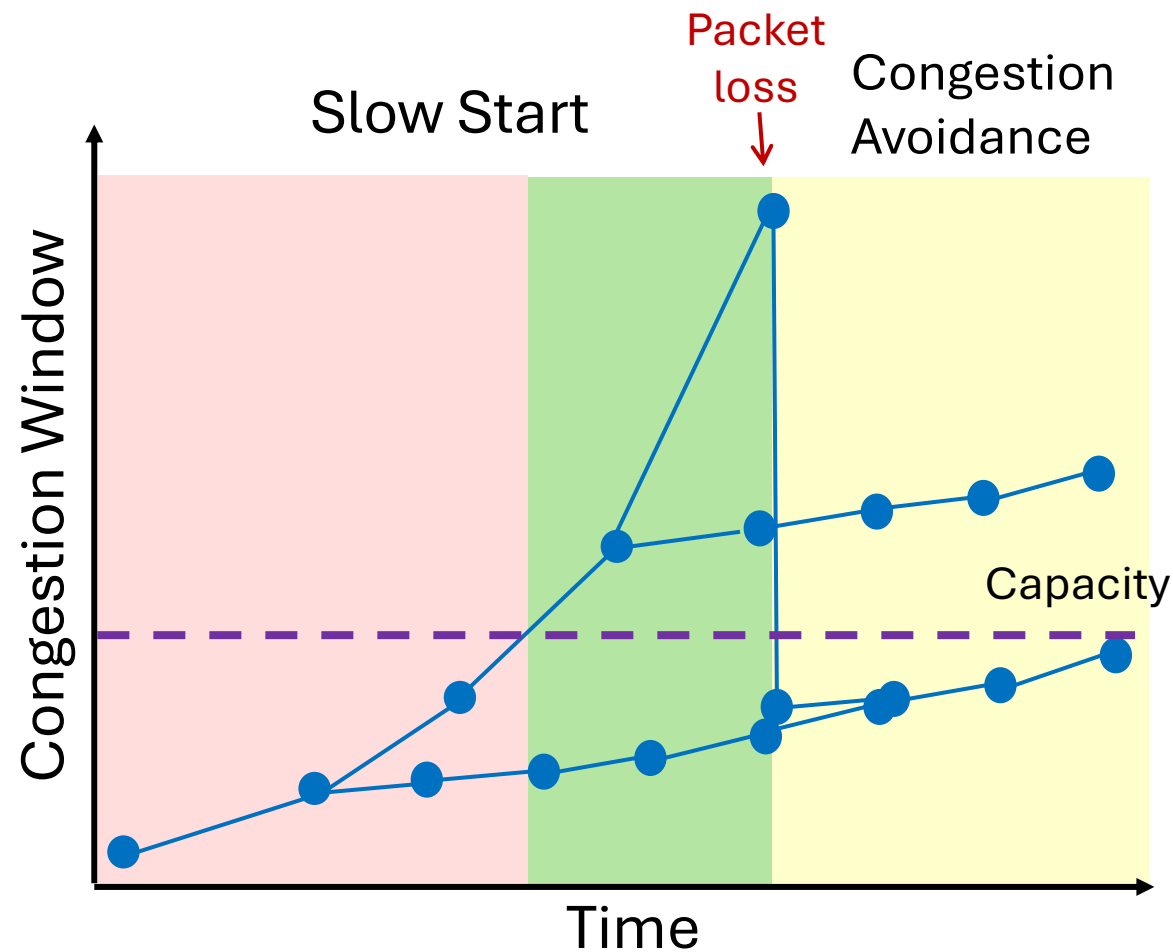
TCP Slow Start

- Congestion window doubles each RTT
- Packet loss → Exit slow start: **Too late (After loss)**
- Idea: Exit after reaching capacity and before loss
- HyStart (Linux default)
- HyStart over wireless

→ **Too early (Before capacity)**

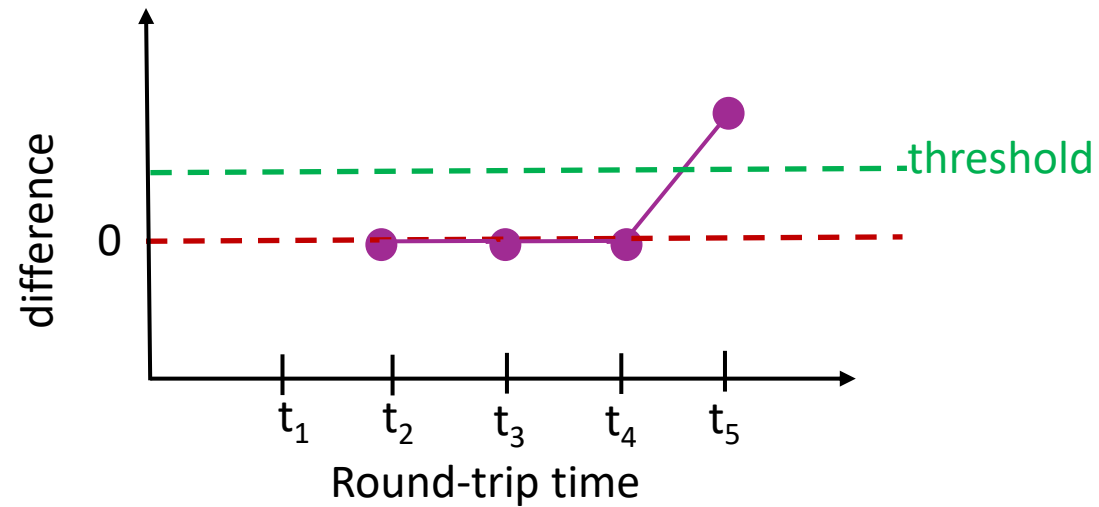
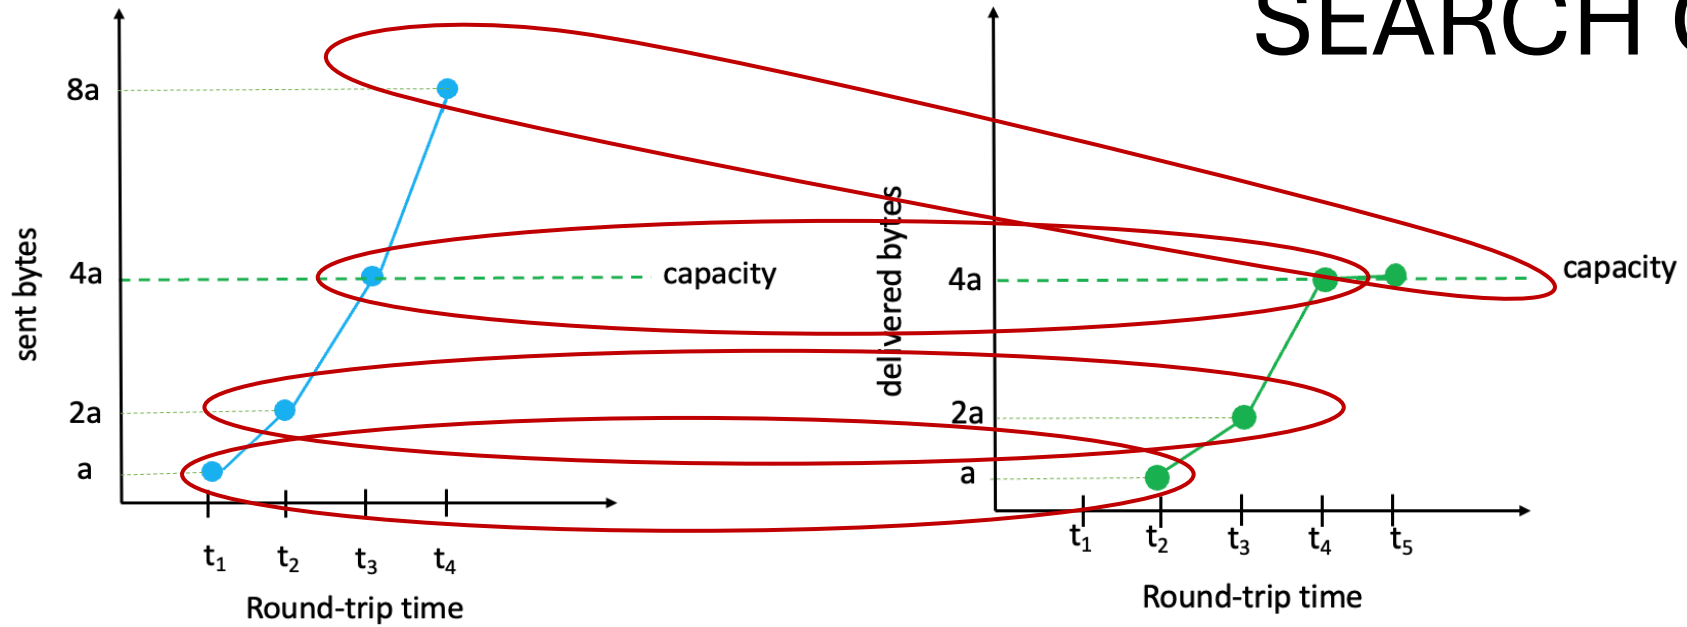
Goal:

→ **At chokepoint: after capacity, before loss**





SEARCH Concept





SEARCH Algorithm

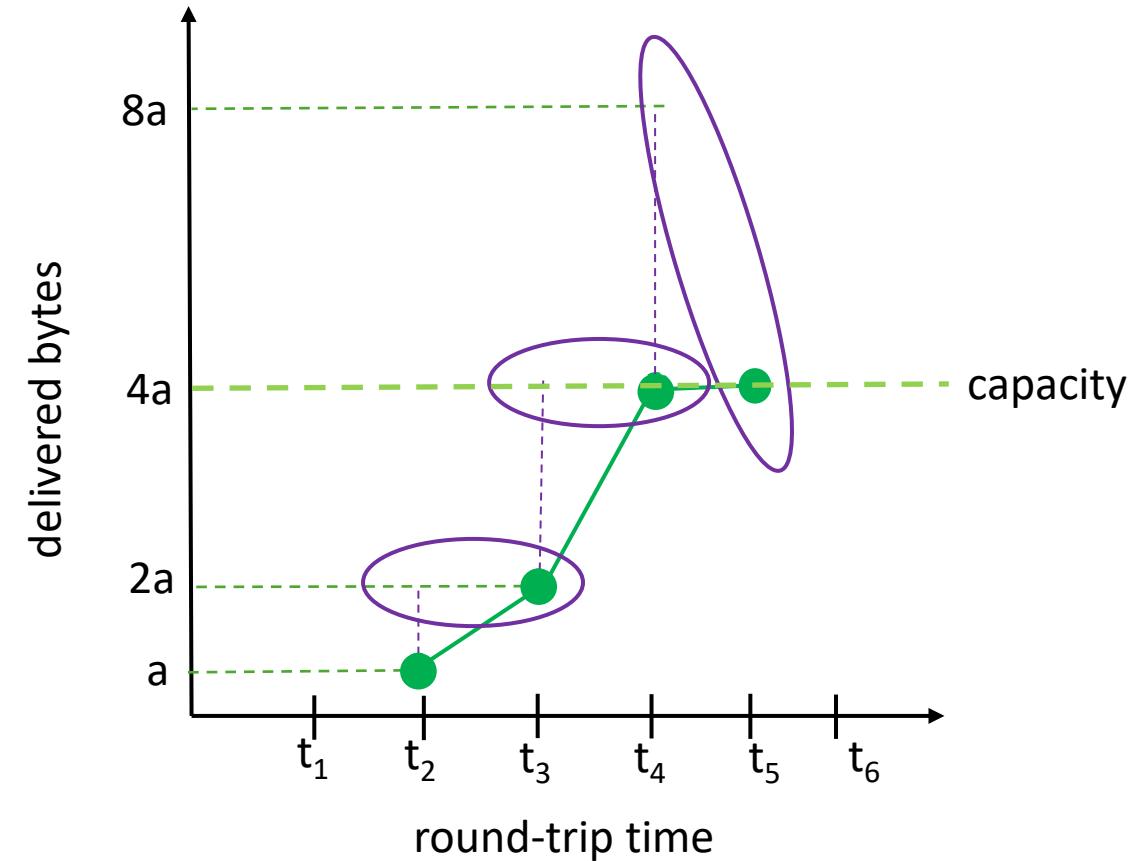
curr_delv

prev_delv

$\text{diff} = 2 \text{ prev_delv} - \text{curr_delv}$

$\text{normalized_diff} = \text{diff} / 2 \text{ prev_delv}$


if ($\text{normalized_diff} \geq \text{threshold}$):
 exit from slow start







Outline

- Place **SEARCH** in slow start space
 - Delay-based: **HyStart**, **HyStart++**
 - Delivered Bytes-based: **BBR**, **SEARCH**
- Memory use delivered byte history
- Test framework
- Summary







Cubic (without HyStart)





Delay-Based: HyStart

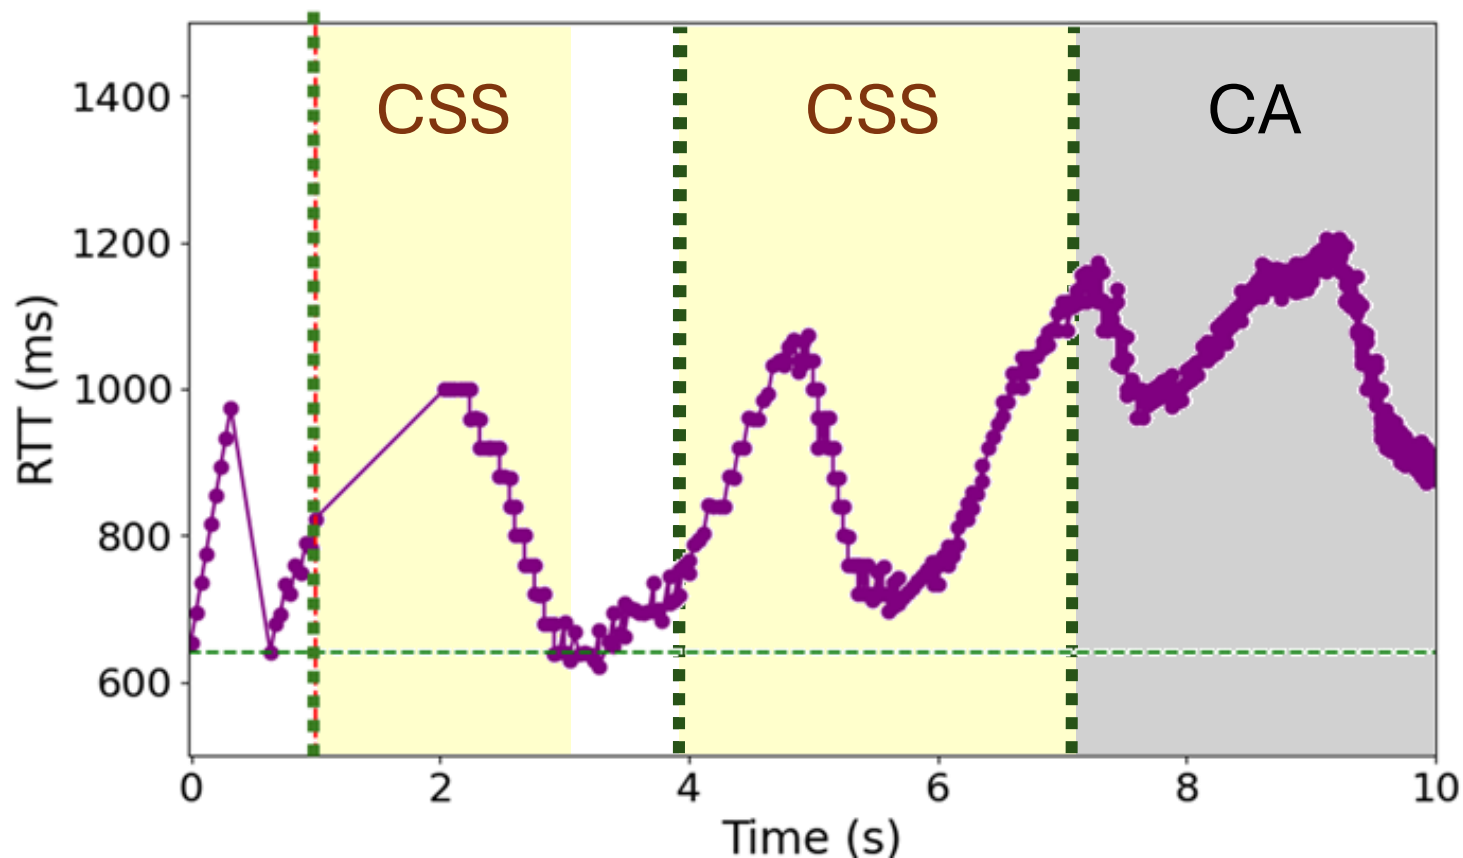
- RTT “spikes” → exit SS
- Sensitive to RTT variation → premature exit, underutilization
 - Bad in high BDP networks





Delay-Based: HyStart++

- Recognizes that delay increase a “noisy” signal
 - Doesn’t have better signal!
- Conservative Slow Start (CSS)
 - Slower growth
 - Check delay
- Delay goes down?
 - Yes → back to SS
 - No → go to CA




- While in CSS grow at slower rate than 2x
- If HyStart was “wrong” – grow slower for some time
- If HyStart was “right” – CSS pushes past capacity limit⁸



Delivered Bytes-Based: BBR



- Exits when delivery rate stops growing
 - 3 consecutive rounds (RTTs), less than 25%
 - Decision each round
- When at capacity, RTTs increase → rounds increase







Delivered Bytes- Based: SEARCH

- Exits when delivered bytes increase less than 30% from previous RTT
- Decision each bin boundary (about $\frac{1}{2}$ RTT)
- Uses initial RTT, not current \rightarrow resilient to congestion delay



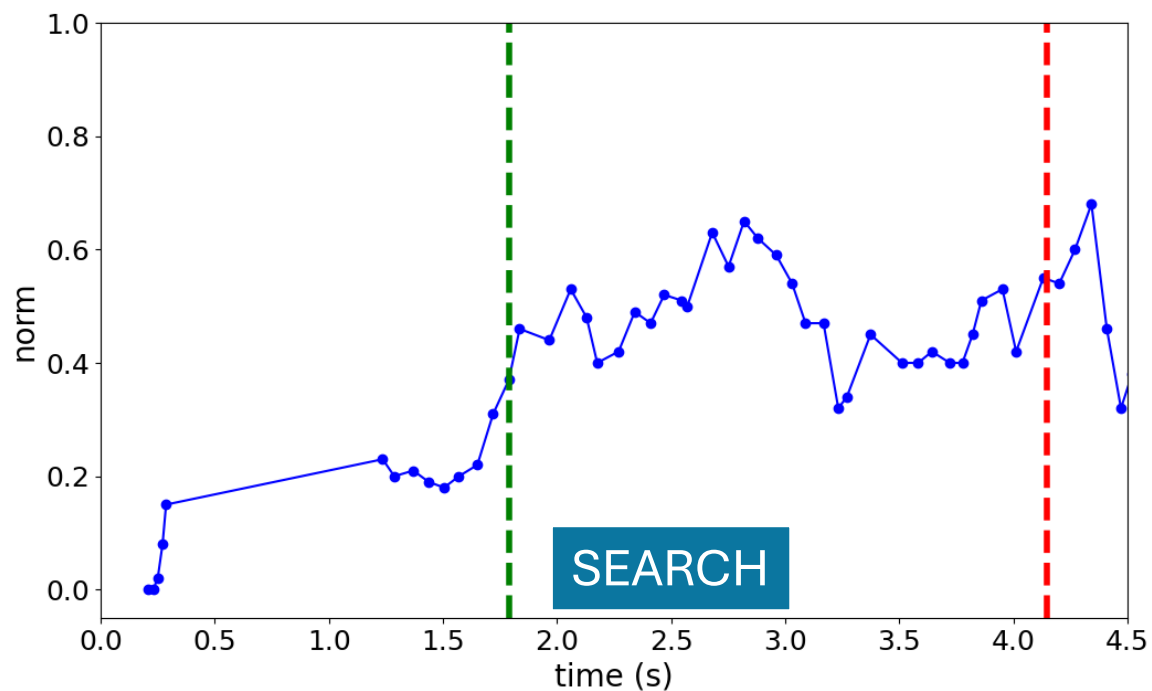
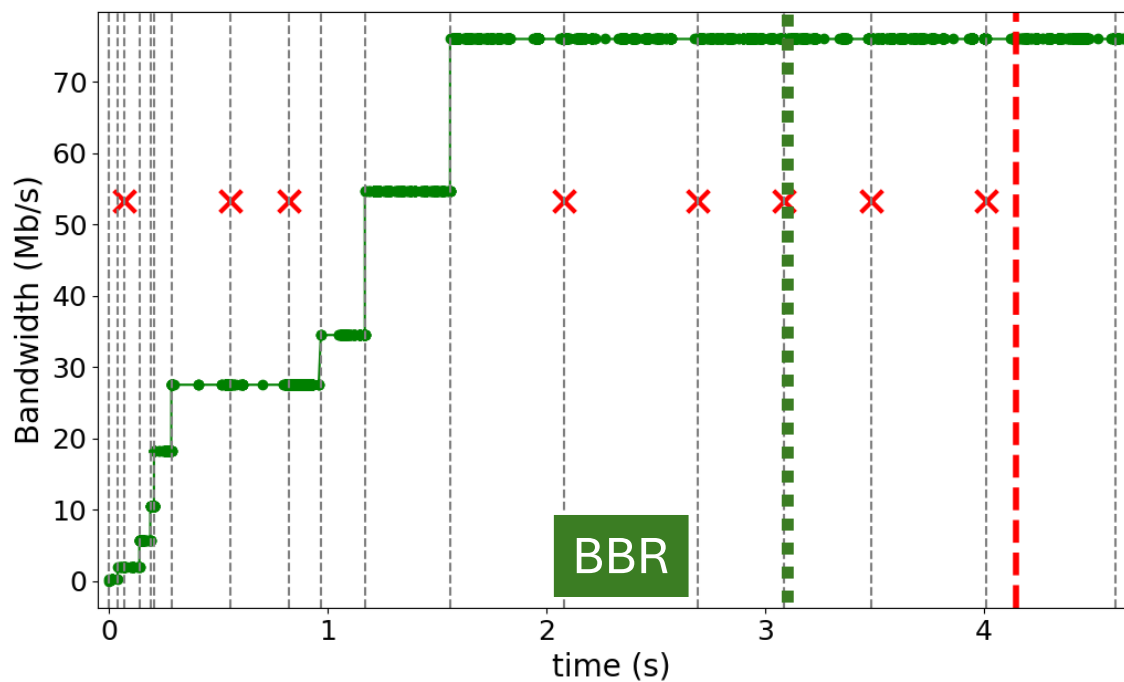
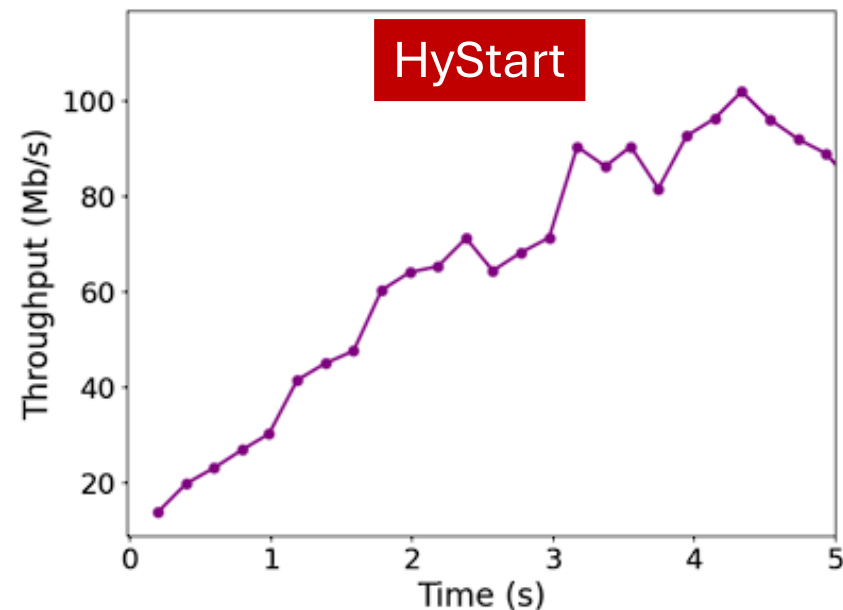
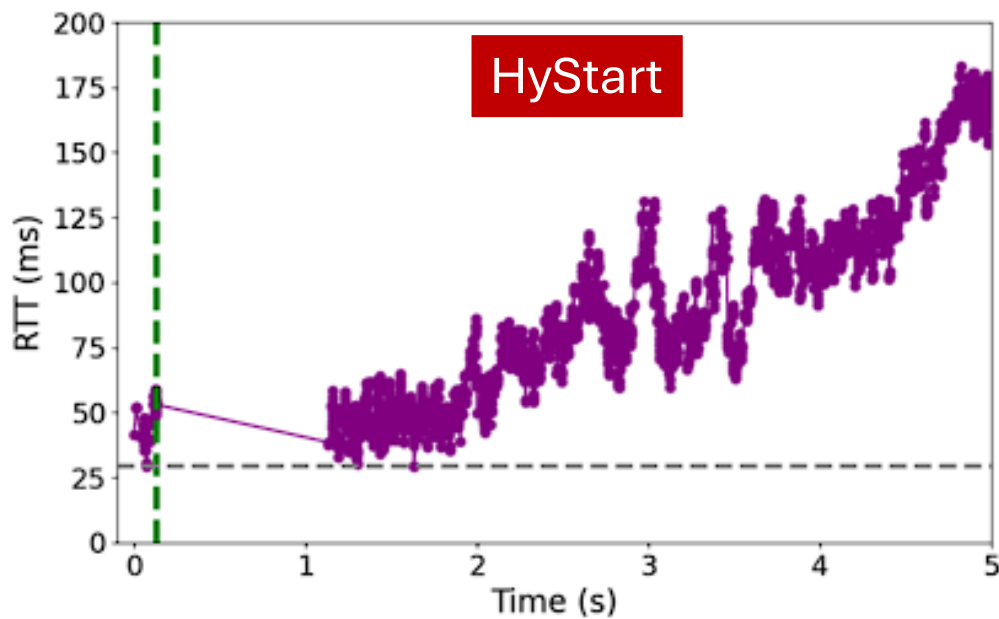


4G LTE – Same old story





4G LTE





Outline

- Place **SEARCH** in slow start space (done)
 - Delay-based: HyStart, HyStart++
 - Delivered Bytes-based: BBR, SEARCH
- Memory use delivered byte history (next)
- Test framework
- Summary




Memory Use for Delivered Bytes

- BBR

- Time when every segment sent
- When segment delivered, compute delivery rate
 - bytes delivered from sent to now / time
- Storage overhead – time size x cwnd size

- SEARCH

- Sequence number each bin ($\sim \frac{1}{2}$ RTT)
- When bin ends, compute delivered
 - seq now – seq then
- Storage overhead – sequence x window




100 MB/s, 100 ms RTT

= 83 total



Window 10 + History 15

= 25 total





Outline

- Place **SEARCH** in slow start space (done)
 - Delay-based: HyStart, HyStart++
 - Delivered Bytes-based: BBR, SEARCH
- Memory use delivered byte history (done)
- Test framework (next)
- Summary



Source Code

Code Extraction


Kernel code
(tcp_cubic_search.c)

```
// SEARCH_defs_begin  
// SEARCH_defs_end  
  
// SEARCH_begin  
// SEARCH_end
```

Challenges:

1. Verify implementations of **SEARCH**
2. Compare performance for parameter tuning


Setup Support Files



Input
CSV file


t_1	s_1
t_2	s_2
...	
t_n	s_n

Input csv



Compilation & Execution

Output file



Output
Log file

Pass
Pass
Fail
Pass



Using the Framework

- Verify and validate **SEARCH** implementations
 - Linux, versions
 - FreeBSD
 - QUIC
- Compare performance
 - Traditional, **HyStart**, **SEARCH**
 - **BBR**, **HyStart++**
- Traces!
 - Normal, edge cases
 - WiFi, 4g/5g, GEO, LEO, Wired


Output
Log file

Pass

Pass

Fail

Pass





Summary



SEARCH

<https://search-ss.wpi.edu/>

- SEARCH

- Determines “choke point” from expected delivered bytes
- Exits slow start after congestion point, before loss

- Framework to verify implementations

- Looking for volunteers to try it out!



claypool@wpi.edu

Thank-you for your attention!

SEARCH – a New Slow Start Algorithm for TCP and QUIC

Jae Chung
Feng Li

Maryam Ataei Kachooei
Mark Claypool

IETF CCWG
Bangkok, Thailand
March 2025





References

- Improving TCP Slow Start Performance in Wireless Networks with SEARCH
 - *IEEE World of Wireless, Mobile and Multimedia Networks (WoWMoM)*
 - Perth, Australia, June 2024
- Improving QUIC Slow Start Behavior in Wireless Networks with SEARCH
 - *IEEE Local and Metropolitan Area Networks (LANMAN)*
 - Boston, Massachusetts, USA, July 2024
- Implementation of the SEARCH Slow Start Algorithm in the Linux Kernel
 - *0x18 NetDev Conference*
 - Santa Clara, California, USA, July 2024
- Reducing Per-flow Memory Use in TCP SEARCH
 - *IEEE World of Wireless, Mobile and Multimedia Networks (WoWMoM)*
 - Fort Worth, TX, USA, May 2025



TCP Congestion Control Test Framework

1. **ss_extract.py** - Extracts Relevant Code

- **Usage:** `ss_extract.py -f original_file.c -k keyword`
- **Functionality:**
 - Extracts slow start-related code using labeled sections (`keyword_begin`, `keyword_end`, `keyword_defs.begin`, `keyword_defs.end`).
 - Generates three key files in `test_dir`:
 1. `keyword_modules.c`
 2. `keyword_defs.h`
 - `tcp.h` (based on extracted modules)



TCP Congestion Control Test Framework

2. `ss_setup.py` - Sets Up Testing Environment

- **Usage:** `ss_setup.py -k keyword`
- **Functionality:**
 - Determines the appropriate test file for the specified protocol.
 - If a predefined `test_keyword.c` exists for the keyword (e.g., SEARCH, HyStart, BBR), it is placed in `test_dir`.
 - If no predefined test file is available, `test_base.c` is provided for user modification.
 - Places `cc_helper_functions.h`, which includes essential congestion control functions, in `test_dir`. Users can extend this file by extracting additional functions from the kernel if needed.



TCP Congestion Control Test Framework

3. `ss_run.py` - Executes the Tests

- **Usage:** `ss_run.py -k keyword -i input -o output`
- **Functionality:**
 - Runs the test with specified input CSV files.
 - Runs `test_keyword`:
 - Reads the input CSV file with test parameters.
 - Applies the congestion control logic.
 - Outputs test results to `output.txt`.



Per-Flow Memory Use

